DATAMARCOS

LIVE ONLINE · MENTOR-LED PROGRAM

AI Security & SecOps

16 weeks · 8–10 hours a week · Part-time, live online

A 16-week, part-time live program for people who want to secure AI and LLM applications, and to use AI carefully inside security operations. You will threat-model and red-team chat, RAG and agent features, design guardrails and basic governance, then practise AI-assisted detection, triage and incident-response automation with a human approval step. Deeper and more hands-on than our 8-week Cybersecurity with AI introduction.

Fees and EMI options shared on a counsellor call.

Program guide · Curriculum, projects and format. No prices are listed in this brochure.

Noida One, 303-304, Rasoolpur Nawada, Industrial Area, Sector 62, Noida, Uttar Pradesh 201309, India

www.datamarcos.com · hello@datamarcos.com · +91 88519 30644

About this program

Who it is for

- Security analysts, SOC, detection and incident-response professionals who want a structured, hands-on path into AI-powered SecOps
- Software, platform and security engineers who ship chat, RAG or agent features and need to secure them in depth
- IT and security leads who need both LLM-application security and a practical view of SIEM/SOAR automation with AI
- Learners who can commit 8–10 hours a week for 16 weeks, joining live from anywhere

Who it is not for

- People looking for an official CompTIA, SANS, GIAC, (ISC)² or government certificate. This program issues a Datamarcos certificate of completion
- Complete beginners who need a first course in cybersecurity fundamentals
- People who only want a short introduction. Our 8-week Cybersecurity with AI course is the lighter starting point
- People seeking placement or salary guarantees — we do not offer those

Starting point

- Some exposure to IT, networks, application logs, security operations or application development is helpful
- No prior machine-learning or LLM-engineering experience is required
- You should be willing to inspect logs, prompts or playbooks and write down what you actually tested
- If you are new to both cybersecurity and AI applications, the 8-week Cybersecurity with AI course is a shorter starting point

What you will be able to do

These are skills and work products from the program — not placement, salary or promotion promises.

- 1. Produce a written threat model of a chat, RAG or agent feature mapped to the OWASP Top 10 for LLM applications
- 2. Run a scoped, documented red-team pass, including prompt injection and data-leakage tests, and record residual risk
- 3. Recommend and sketch guardrails: input and output review, least privilege for tools, fallbacks and a human approval step
- 4. Draft basic AI-governance artefacts: an acceptable-use note, a logging and review plan, and when to involve legal or compliance review
- 5. Design an AI-assisted detection or triage workflow in a SIEM-style setting, with every model claim verified against source evidence
- 6. Build a SOAR-style response playbook with approvals, least privilege, verification and a kill switch — not unsighted automation
- 7. Leave with a capstone pack: an LLM security assessment and a SecOps playbook, not a vendor or professional pentest certification

Projects

Project 1: LLM application threat model and red-team note

A written threat model mapped to the OWASP Top 10 for LLM applications, plus a scoped, documented red-team pass with residual risk.

Project 2: Guardrails and governance pack

Input and output review, least-privilege tool use, a logging and review plan, and a short acceptable-use note. Not a legal or official compliance certificate.

Project 3: Capstone: AI-assisted SecOps playbook

A verified triage or detection workflow and a SOAR-style response playbook with approvals, verification and a kill switch, presented with the LLM security assessment.

Curriculum · 16 weeks

8–10 hours a week, live and mentor-led. Open each week in the live sessions for labs and feedback.

Week 1: The AI and LLM attack surface

- How chat, retrieval and tool-using agents change the attack surface
- What stays the same as traditional application security
- Assets, trust boundaries and where data flows
- Picking one feature you will carry through the first half of the program

Lab: Sketch the components and data flows of a sample LLM feature.

Week 2: OWASP Top 10 for LLM applications, in depth

- Walking the OWASP Top 10 for LLM applications as a working checklist, not as a certificate
- Prompt injection, data leakage, insecure output handling and supply-chain issues, in practical terms
- Agent and tool-use risks: least privilege and unexpected actions
- How this list should and should not be used in a review

Lab: Score a sample application against the list and note gaps in evidence.
You ship: An OWASP-LLM worksheet with evidence notes.

Week 3: Threat modelling an LLM, RAG or agent feature

- A lightweight method: assets, threats, existing controls, residual risk
- Indirect prompt injection via retrieved documents or tool output
- What you would need from engineering to test each item
- Writing the model so a non-specialist can read it

Lab: Draft a threat model for your chosen feature.
You ship: A first-pass threat model.

Week 4: Phase checkpoint: a review-ready threat model

- Tightening evidence, residual risk and recommended tests
- Peer and mentor review
- What the hands-on security labs will add next

Lab: Present the threat model and the tests you plan to run.
You ship: A review-ready threat model and a test backlog.

Week 5: Prompt injection, in the lab

- Direct and indirect prompt injection, including documents and tool output
- Jailbreak-style attempts used only on a practice application you are authorised to test
- Recording steps so an engineer can reproduce a finding
- What a single successful injection does and does not prove

Lab: Run an authorised injection test plan on a practice app and log results.
You ship: An injection findings log.

Curriculum (continued)

Week 6: Data leakage, output handling and tool misuse

- Sensitive data in prompts, context windows, logs and retrieved documents
- When a model reply is treated as a trusted instruction or as code
- Tool and agent misuse: unexpected actions, over-broad permissions
- Designing a least-privilege tool interface

Lab: Attempt leakage and insecure-output cases, then tighten a tool permission in the practice app.
You ship: A leakage and tool-misuse note with suggested controls.

Week 7: Guardrails you can actually run

- Input and output review, allow-lists, structured outputs and fallbacks
- Human-in-the-loop steps that will still be followed under time pressure
- Evaluation gates before a prompt or a tool change is released
- Limits of filters: they reduce risk, they do not eliminate it

Lab: Add or document a small set of guardrails and re-test an earlier finding.
You ship: A guardrails design and a re-test note.

Week 8: A scoped red-team pass

- Rules of engagement, scope and what is out of bounds
- A short test plan mapped to the threat model
- Residual risk: what you did not test
- This is structured coursework, not a professional penetration-testing or official red-team certification

Lab: Execute the test plan and write findings an engineer can act on.
You ship: A red-team report with residual risk.

Week 9: AI governance basics for teams

- Acceptable use, ownership and residual-risk language that does not over-claim
- Public frameworks such as the NIST AI Risk Management Framework used as a map, not as a certificate
- When to involve legal, privacy or compliance review
- A one-page governance note a team can actually adopt

Lab: Draft an acceptable-use and ownership note for your feature.
You ship: A short governance note.

Week 10: Logging, monitoring and incident notes for AI systems

- What to log (prompts, retrieval, tool calls) and what must be redacted
- Alerts that are useful versus noise
- An incident note when an LLM feature misbehaves or is abused
- Handover into a wider security-operations process

Lab: Write a logging and review plan and a sample incident note.

Curriculum (continued)

Week 11: SIEM, SOAR and where AI should sit

- What a SIEM and a SOAR-style platform are for, in practical terms
- Where AI can help: summarising, clustering, suggesting queries or next steps
- What must stay a human decision: containment, customer notification, legal holds
- You will practise in a SIEM or SOAR-style workflow you can access, or in a practice environment — Datamarcos is not a vendor for those products

Lab: Map a current or sample alert path and mark where AI may and may not be used.
You ship: A SecOps-plus-AI system sketch.

Week 12: Detection and log analysis with AI, verified

- Writing precise questions over logs or alerts so answers can be checked
- Drafting or translating detection logic with an assistant, then testing it
- Verifying indicators, timestamps and asset names against the source
- A method note: what the model suggested and what you confirmed

Lab: Run an assisted analysis on a sample log or alert set and tick off every claim.
You ship: A verified detection or log-analysis note.

Week 13: Alert triage and enrichment

- Enrichment, clustering and a first-pass severity suggestion
- True-positive and false-positive discipline
- Writing a triage brief another analyst can pick up
- What not to paste into a public assistant

Lab: Triage a small alert batch with AI assistance and a review checklist.
You ship: A triage pack with enrichment notes and review points.

Week 14: Incident response automation with a kill switch

- Turning a playbook into a SOAR-style workflow: triggers, conditions, actions
- Approvals, least privilege, verification and recovery
- Kill switches and rollback when automation is wrong
- Auto-isolation and similar actions only with a clear approval path

Lab: Design a playbook for one incident type, including the approval and stop conditions.
You ship: A SOAR-style playbook with human gates.

Week 15: Metrics, failure testing and when automation hurts

- What to measure: time to triage, quality of summaries, automation error rate
- Testing a playbook against a case where the obvious action would be wrong
- Staged rollout of automation, not a single big-bang switch
- A written recommendation on what to automate first, and what to leave manual

Lab: Failure-test your playbook and revise the gates.

Curriculum (continued)

Week 16: Capstone: assessment, playbook and walk-through

- Assembling the threat model, red-team note, guardrails and governance artefacts
- Assembling the verified SecOps workflow and the SOAR-style playbook
- Live presentations and mentor feedback
- Honest scope: coursework, not an official pentest, GIAC, SANS or vendor certificate

Lab: Present the capstone pack and the checks you ran.

Format, fees and refunds

Part-time and live online, in a small mentor-led cohort. Weekend-morning batches are offered for learners in the US. Evening batches (India time) are offered for learners in the Gulf and Africa. Europe-oriented groups typically meet after work and on Saturdays. Recordings are shared if you miss a live session. Exact days, times and the next cohort date are confirmed on a counsellor call.

Join live from anywhere. Weekend-morning batches are offered for learners in the US; evening batches (India time) for the Gulf and Africa; Europe-oriented groups typically meet after work and on Saturdays.

A Datamarcos certificate of completion is awarded when you finish the program work. Cohort start dates are shared on a counsellor call — we do not publish a date until then.

Fees and EMI options shared on a counsellor call.

7-day refund policy: if the program is not right for you, ask within 7 days of enrollment.

Mentor

Live sessions are mentor-led by Deepesh Mishra, founder of Datamarcos, a BITS Pilani alumnus, and a solution architect and corporate trainer in data and AI who has trained teams at 200+ organizations.

Questions, answered

How is this different from Cybersecurity with AI?

Cybersecurity with AI is an 8-week introduction: a first pass at AI in security operations and at securing LLM applications. AI Security & SecOps is a 16-week, more hands-on program. It goes further on threat modelling, prompt injection labs, guardrails, basic AI governance, and SIEM/SOAR-style detection, triage and incident-response automation with approvals and a kill switch. If you want a shorter first course, start with Cybersecurity with AI.

Is this an official CompTIA, SANS, GIAC or government certificate?

No. Completing the program leads to a Datamarcos certificate of completion. We use public frameworks such as the OWASP Top 10 for LLM applications, and we refer to maps such as the NIST AI Risk Management Framework, as working checklists. We do not award, and do not prepare you for, a third-party certification exam.

How long is the program and how much time does it take?

Sixteen weeks, part-time and live online, at about 8–10 hours a week including the live session and practice work.

Who is it for?

Security, SOC, detection and incident-response professionals, and engineers who ship LLM or agent features and need to secure them in depth. Some IT, logging or security exposure is helpful. It is not a first course in cybersecurity fundamentals.

Are the sessions live?

Yes. Sessions are live and mentor-led, with recordings if you miss a class. Weekend-morning batches are offered for learners in the US, with other regional times available. Exact times are confirmed on a counsellor call.

Will you teach a specific SIEM or SOAR product?

You will practise SIEM- and SOAR-style workflows — detection, triage, enrichment and playbooks with human gates — in a tool you can access or in a practice environment. Datamarcos is not a vendor for those products, and the program is not a vendor certification.

What about fees?

Fees and EMI options shared on a counsellor call.

Is there a refund policy?

Yes. Datamarcos has a 7-day refund policy. If the program is not right for you, ask within 7 days of enrollment.

DATAMARCOS

Talk to a counsellor

Share a few details on datamarcos.com and a counsellor will talk through program fit, the next cohort, fees and enrollment.

Fees and EMI options shared on a counsellor call.

7-day refund policy: if the program is not right for you, ask within 7 days of enrollment.

Noida One, 303-304, Rasoolpur Nawada, Industrial Area, Sector 62, Noida, Uttar Pradesh 201309, India

www.datamarcos.com · hello@datamarcos.com · +91 88519 30644

This brochure does not promise placement, salary or promotion outcomes. It is not legal, immigration or official-government certification advice.