



2-MONTH CERTIFICATION PROGRAM

APPLIED AGENTIC AI CERTIFICATION

Go beyond using ChatGPT, Claude, or Copilot. Learn how Agentic AI systems are actually designed and built — agent architecture, tool calling, APIs, agentic workflows, and the frameworks powering production AI today.

8-Week Practical Program

n8n · LangChain · LangGraph · MCP

Live Capstone Project

Mentor-led by Deepesh Mishra

Just a short while ago, OpenAI announced a new model, GPT-6 Astra.

And this is what most professionals are missing: until now, even some of the strongest AI models were scoring around **30%** on ARC-AGI-3 — a benchmark that tests how well AI can solve completely new problems. With Astra, early reports suggest scores close to **99.9%**.

We're getting closer to that Jarvis-like world where AI doesn't just answer questions — it can reason, use tools, work with software, and actually take action. This is exactly why the conversation around **AGI (Artificial General Intelligence)** is heating up. In simple words: AI that can handle a wide range of intellectual tasks beyond the human level.

One question comes to mind — have we reached AGI? No. One benchmark doesn't prove that. But the direction is crystal clear.

Six months from now, many professionals may look back at September 2026 and think, *"I should have started learning this technology then."* And you don't need to wait for AGI to prepare. The practical shift is already here: **Agentic AI**.

KNOW YOUR MENTOR



Deepesh Mishra

Founder, Datamarcos · Solution Architect · Corporate Trainer

"Hey, I am Deepesh Mishra. I am a Solution Architect and Corporate Trainer with over 18 years of industry experience. I have trained more than 25,000 professionals across 200+ organisations, including professionals and teams associated with Deloitte, EY, Adobe, HSBC, Boeing, and Reliance.

In my 2-month Applied Agentic AI Certification, we go beyond simply using ChatGPT, Claude, or Copilot. We focus on how Agentic AI systems are designed and built — with practical work around agent architecture, tool calling, APIs, agentic workflows, and relevant frameworks such as n8n, LangChain, LangGraph, MCP, and automation platforms where appropriate."

25,000+

PROFESSIONALS TRAINED

200+

ORGANIZATIONS

18+

YEARS EXPERIENCE

10+

COUNTRIES

Organizations Whose Teams Deepesh Has Trained

Deloitte

EY

Adobe

HSBC

Boeing

Reliance

PwC

Bank of America

Yahoo Singapore

and 200+ more

Build Systems That Reason, Plan, and Take Action

Progress from foundational LLM concepts to deploying real, industry-grade AI agents in just 8 weeks. Through hands-on, practical training, you'll design and build a multi-SDK LLM playground, production RAG pipelines, tool-using agents, multi-agent systems, and a fully integrated capstone application.

**8-Week
Program**

**1:1
Mentorship**

**Live
Projects**

**Career
Support**

Who Is This Course For

- ✓ Automation QA / Test Automation Professionals
- ✓ Software Engineers & Developers (Backend, Frontend, Full Stack)
- ✓ Technical Product Managers
- ✓ Professionals exploring AI-powered development workflows
- ✓ Early-career professionals building practical AI skills

Prerequisites & Eligibility

You don't need to be an AI expert or come from a technical background to join.

- ✓ Open to learning basic Python programming
- ✓ Comfortable exploring new tools and hands-on exercises
- ✓ No prior AI experience required — beginner-friendly refreshers provided before live sessions

From LLM Foundations to Reliable, Structured AI Systems

Week 1 · LLM Foundations & System Thinking

FOUNDATIONS

How LLMs actually work: transformer intuition, attention, tokenization & embeddings, context window limits, hallucination causes, and prompt vs. retrieval vs. fine-tuning — enough to engineer real systems around them. Hands-on with the OpenAI, Gemini & Claude SDKs.

Transformers & Attention

Tokens & Embeddings

Multi-SDK Basics

Project: Multi-SDK LLM Playground — query OpenAI, Gemini & Claude side by side

Week 2 · Prompt Engineering & Structured Output

PROMPTING

Zero-shot, few-shot, chain-of-thought & tree-of-thought prompting, role/persona design, structured JSON output, function calling, and defending against prompt injection — with an evaluation-driven test-measure-improve loop.

Advanced Prompting

Function Calling

Prompt Injection Defense

Project: Call Center Transcript Sentiment & Intent Classifier with an evaluation harness

Week 3 · Production RAG Systems

RETRIEVAL

Why RAG is essential for enterprise AI, end-to-end RAG architecture, embeddings, keyword vs. semantic search, vector databases, chunking strategies, retrieve-then-rerank pipelines, and evaluating faithfulness & relevance with RAGAS.

Vector Databases

Chunking Strategies

RAGAS Evaluation

Project: Company Knowledge-Base Chatbot — ingest PDFs & docs with hybrid search + reranking

Week 4 · Building AI Agents

AGENTS

Core components of an AI agent — roles, memory, tools, autonomy, goals. The ReAct framework, tool orchestration, short/long-term memory systems, and connecting agents to APIs, databases and file systems safely.

ReAct Framework

Agent Memory

Tool Orchestration

Project: Researcher-and-Summarizer Agent + a Customer Care Agent using live tool calling

From Single Agents to Production Multi-Agent Systems

Week 5 · Multi-Agent Systems & Orchestration

ORCHESTRATION

Multi-agent architectures — sequential, parallel, hierarchical. Role-based specialist vs. generalist agents, coordinator-dispatcher patterns, and passing context/outputs across agent boundaries at production scale.

CrewAI

LangGraph

Coordinator-Dispatcher Design

Project: *Financial News Research Coordinator — dispatches specialist sub-agents & produces a structured report*

Week 6 · Model Context Protocol (MCP) & Enterprise Integration

MCP

What MCP is and why modern agents use it, client-server-transport architecture, registering tools/services through MCP servers, and how it simplifies connecting agents to databases, APIs and internal systems at scale.

MCP Architecture

Enterprise Tool Access

Migration from Legacy Tool-Calling

Project: *MCP server connecting an agent to a live database and internal API*

Week 7 · Voice, Multimodal Agents & Automation

MULTIMODAL

Multimodal capabilities across vision, audio and documents, voice agent architecture (speech-to-text → LLM → action → text-to-speech), and wiring agentic workflows into no-code automation platforms like n8n.

Voice Agents

Multimodal AI

n8n Automation

Project: *Voice-enabled multimodal assistant with an automated workflow trigger*

Week 8 · Evaluation, Reliability & Capstone

CAPSTONE

Evaluating agents end-to-end (task completion, tool-use accuracy, response quality), debugging agentic systems, reliability patterns (idempotency, retries, fallbacks), and shipping a fully integrated capstone.

Agent Evaluation

Reliability Patterns

Capstone Deployment

Capstone: *Fully integrated multi-agent application combining RAG + Agents + Multi-Agent + MCP*

EVERY TOOL YOU WILL LEARN AND USE

Technologies & Frameworks You Will Master

Hands-on with the exact SDKs, frameworks and platforms production Agentic AI teams use today
— not just theory slides.

Python	OpenAI SDK	Gemini SDK	Claude SDK	tiktoken
LangChain	LangGraph	CrewAI	Google ADK	n8n
FAISS	ChromaDB	Pinecone	Weaviate	RAGAS
LangSmith	Langfuse	Guardrails AI	MCP SDK	FastAPI

What You'll Build (Live Projects)

- ✓ Multi-SDK LLM Playground
- ✓ Sentiment & Intent Classifier
- ✓ Production RAG Knowledge-Base Chatbot
- ✓ Researcher & Customer Care Agents
- ✓ Multi-Agent Research Coordinator
- ✓ Capstone: Integrated Agentic AI Application

Why This, Why Now

- ✓ AI is moving from "answering questions" to reasoning, using tools, and taking action
- ✓ Companies are hiring for engineers who can build agentic systems, not just prompt a chatbot
- ✓ The practical shift toward Agentic AI is already here — you don't need to wait for AGI to prepare

Real Experience. Real Impact. Real Growth.

Academic Excellence

BITS Pilani Alumni — strong foundation, proven excellence.

AI Tool Power User

Hands-on expert across leading AI platforms and modern agentic stacks.

18+ Years of Experience

Corporate Trainer, Solution Architect, Tech Expert.

Industry Trusted Trainer

Worked with global leaders like Deloitte, EY, Adobe, HSBC, Boeing & Reliance.

25,000+

PROFESSIONALS TRAINED

200+

CORPORATE CLIENTS

10+

COUNTRIES

1,000+

SESSIONS DELIVERED

"His deep understanding of the subject, combined with a clear and engaging teaching style, made complex concepts easy to grasp — even for those new to the field."

Kumar Gaurav — Senior Azure Data Engineer @ PwC, Ex-EY, Ex-Genpact

Practical. Relevant. Impactful.

No theory-only sessions. Every concept is taught through hands-on labs and live projects that mirror what you'll actually build on the job.

Live, Mentor-Led Sessions

Small-batch, interactive sessions led directly by Deepesh Mishra, with real-time Q&A and doubt resolution.

Hands-on Labs & Projects Every Week

Every module includes real lab work — you build the agent, not just watch a demo.

Fully Integrated Capstone

End with a portfolio-ready capstone combining RAG, Agents, Multi-Agent orchestration and MCP.

Continued Mentorship

Support doesn't end at the classroom — get guidance beyond sessions as you apply these skills at work.

Program Includes

- ✓ Hands-on projects every week
- ✓ Industry-relevant, up-to-date curriculum
- ✓ Learn directly from an industry expert
- ✓ Capstone project for job-readiness
- ✓ High-growth career opportunities
- ✓ Certificate of Completion

Where Applied Agentic AI Takes You

Whether you're in QA automation, backend development, or product management — Agentic AI skills are becoming the differentiator for the next wave of hiring.

Starting Point	What Changes After This Program
Automation QA / Test Engineer	Builds and evaluates AI agents, not just test scripts
Software / Backend Developer	Designs tool-using, agentic systems on top of existing APIs
Technical Product Manager	Speaks fluently on agent architecture and can scope AI-native products
Data / AI Professional	Moves from "using AI tools" to engineering production agentic systems

Roles You Become Ready For

AI Agent Engineer

Agentic AI Developer

GenAI Solutions Engineer

AI Automation Engineer

AI Product Engineer

Certificate of Completion

Every learner who completes the program receives a Datamarcos Applied Agentic AI Certification, validating hands-on proficiency across LLMs, RAG, agents, multi-agent orchestration and MCP.

GET STARTED

Six Months From Now, You Could Be Shipping Agents.

You don't need to wait for AGI to prepare. The practical shift is already here. Reach out to the Datamarcos team to know upcoming batch dates, fees, and enrollment steps for the Applied Agentic AI Certification.

Ready to Build What's Next?

Become an Applied Agentic AI Engineer in 8 weeks.

CONTACT US TO ENROLL

OFFICE ADDRESS

C-908, Sector 62, Noida, Uttar Pradesh, India

WEBSITE

www.datamarcos.com

FOUNDER & MENTOR

Deepesh Mishra

PROGRAM

Applied Agentic AI Certification (2 Months / 8 Weeks)



DATAMARCOS

© Datamarcos IT Consulting Firm · All Rights Reserved